

CS145 Howework 1

****Important Note:**** HW1 is due on **11:59 PM PT, Oct 19 (Monday, Week 3)**. Please submit through GradeScope (you will receive an invite to Gradescope for CS145 Fall 2020.).

Print Out Your Name and UID

****Name:** Ali Mirabzadeh , **UID:** 305179067

Before You Start

You need to first create HW1 conda environment by the given `cs145hw1.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda create -f cs145hw1.yml
conda activate hw1
conda deactivate
```

More useful information about managing environments can be found [here](https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html) (<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks that you are allowed to edit (between `START/END YOUR CODE HERE`), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [1]: import numpy as np
import pandas as pd
import sys
import random as rd
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
%load_ext autoreload
%autoreload 2
```

If you can successfully run the code above, there will be no problem for environment setting.

1. Linear regression

This workbook will walk you through a linear regression example.

```
In [2]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 100) and training labels (1000,)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)
```

```
Training data shape: (1000, 100)
```

```
Training labels shape: (1000,)
```

1.1 Closed form solution

In this section, complete the `getBeta` function in `linear_regression.py` which use the close for solution of $\hat{\beta}$.

Train you model by using `lm.train('0')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [3]: from hw1code.linear_regression import LinearRegression

lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####
beta = lm.train('0')
train_predict = lm.predict(lm.train_x.values, beta)
test_predict = lm.predict(lm.test_x.values, beta)
training_error = lm.compute_mse(train_predict, lm.train_y.values)
testing_error = lm.compute_mse(test_predict, lm.test_y.values)
#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)

## below for normalizing
lm.normalize()
beta_normalized = lm.train('0')
train_predict_n = lm.predict(lm.train_x.values, beta_normalized)
test_predict_n = lm.predict(lm.test_x.values, beta_normalized)
training_error_n = lm.compute_mse(train_predict_n, lm.train_y.values)
testing_error_n = lm.compute_mse(test_predict_n, lm.test_y.values)
print('Normalized training error is: ', training_error_n)
print('Normalized testing error is: ', testing_error_n)

Learning Algorithm Type: 0
Training error is: 0.08693886675396784
Testing error is: 0.11017540281675804
Learning Algorithm Type: 0
Normalized training error is: 0.08693886675396784
Normalized testing error is: 0.11017540281675804
```

1.2 Batch gradient descent

In this section, complete the `getBetaBatchGradient` function in `linear_regression.py` which compute the gradient of the objective fuction.

Train you model by using `lm.train('1')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```
In [4]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####
beta = lm.train('1')
train_predict = lm.predict(lm.train_x.values, beta)
test_predict = lm.predict(lm.test_x.values, beta)
training_error = lm.compute_mse(train_predict, lm.train_y.values)
testing_error = lm.compute_mse(test_predict, lm.test_y.values)
#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)

## below for normalizing
lm.normalize()
beta_normalized = lm.train('1')
train_predict_n = lm.predict(lm.train_x.values, beta_normalized)
test_predict_n = lm.predict(lm.test_x.values, beta_normalized)
training_error_n = lm.compute_mse(train_predict_n, lm.train_y.values)
testing_error_n = lm.compute_mse(test_predict_n, lm.test_y.values)
print('Normalized training error is: ', training_error_n)
print('Normalized testing error is: ', testing_error_n)
```

```
Learning Algorithm Type: 1
Training error is: 0.08693919081192665
Testing error is: 0.11019432186477264
Learning Algorithm Type: 1
Normalized training error is: 0.09654351096262996
Normalized testing error is: 0.12865993449113894
```

1.3 Stochastic gadient descent

In this section, complete the `getBetaStochasticGradient` function in `linear_regression.py`, which use an estimated gradient of the objective function.

Train you model by using `lm.train('2')` function.

Print the training error and the testing error using `lm.predict` and `lm.compute_mse` given.

```

In [5]: lm=LinearRegression()
lm.load_data('./data/linear-regression-train.csv', './data/linear-regression-test.csv')
training_error= 0
testing_error= 0
#####
# STRART YOUR CODE HERE #
#####
beta = lm.train('2')
train_predict = lm.predict(lm.train_x.values, beta)
test_predict = lm.predict(lm.test_x.values, beta)
training_error = lm.compute_mse(train_predict, lm.train_y.values)
testing_error = lm.compute_mse(test_predict, lm.test_y.values)
#####
# END YOUR CODE HERE #
#####
print('Training error is: ', training_error)
print('Testing error is: ', testing_error)

## below for normalizing
lm.normalize()
beta_normalized = lm.train('2')
train_predict_n = lm.predict(lm.train_x.values, beta_normalized)
test_predict_n = lm.predict(lm.test_x.values, beta_normalized)
training_error_n = lm.compute_mse(train_predict_n, lm.train_y.values)
testing_error_n = lm.compute_mse(test_predict_n, lm.test_y.values)
print('Normalized training error is: ', training_error_n)
print('Normalized testing error is: ', testing_error_n)

Learning Algorithm Type: 2
Training error is: 91.10508480511325
Testing error is: 96.50500424208748
Learning Algorithm Type: 2
Normalized training error is: 31.219733063437648
Normalized testing error is: 29.420724030847268

```

Questions:

1. Compare the MSE on the testing dataset for each version. Are they the same? Why or why not?
2. Apply z-score normalization for eachh featurre and comment whether or not it affect the three algorithm.
3. Ridge regression is adding an L2 regularization term to the original objective function of mean squared error. The objective function become following:

$$J(\beta) = \frac{1}{2n} \sum_i (x_i^T \beta - y_i)^2 + \frac{\lambda}{2n} \sum_j \beta_j^2,$$

where $\lambda \geq 0$, which is a hyper parameter that controls the trade off. Take the derivative of this provided objective function and derive the closed form solution for β .

Your answer here:

1. With some small precision difference, they are the same. We see the error rates to vary because these are different models. The stochastic gradient descent has the largest error rate because our dataset is not large enough for this method to be efficient. More precise methods like batch gradient descent and linear regression are more effective.
2. You can see each section (1.1, 1.2, 1.3) as I have written the code for it and see the printed result. With the normalized data we can see an increase in error rate for batch gradient descent, no change in error rate for linear regression and a decrease in error rate for stochastic gradient descent.
- 3.

$$\beta = (x^T x + \lambda)^{-1} x^T y$$

2. Logistic regression

This workbook will walk you through a logistic regression example.

```
In [6]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
# As a sanity check, we print out the size of the training data (1000, 5) and training labels (1000,)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)

Training data shape: (1000, 5)
Training labels shape: (1000,)
```

2.1 Batch gradient descent

In this section, complete the `getBeta_BatchGradient` in `logistic_regression.py`, which compute the gradient of the log likelihood function.

Complete the `compute_avglogL` function in `logistic_regression.py` for sanity check.

Train your model by using `lm.train('0')` function.

And print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```
In [7]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####
lm.normalize()
beta = lm.train('0')
train_predict = lm.predict(lm.train_x.values, beta)
test_predict = lm.predict(lm.test_x.values, beta)
training_accuracy = lm.compute_accuracy(train_predict, lm.train_y.values)
testing_accuracy = lm.compute_accuracy(test_predict, lm.test_y.values)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)
```

```
average logL for iteration 0: -0.5561796575778383
average logL for iteration 1000: -0.4615656578257467
average logL for iteration 2000: -0.4615656578257467
average logL for iteration 3000: -0.4615656578257467
average logL for iteration 4000: -0.4615656578257467
average logL for iteration 5000: -0.4615656578257467
average logL for iteration 6000: -0.4615656578257467
average logL for iteration 7000: -0.4615656578257467
average logL for iteration 8000: -0.4615656578257467
average logL for iteration 9000: -0.4615656578257467
Training avgLogL: -0.4615656578257467
Training accuracy is: 0.798
Testing accuracy is: 0.7495029821073559
```

2.2 Newton Raphhson

In this section, complete the `getBeta_Newton` in `logistic_regression.py`, which make use of both first and second derivative.

Train you model by using `lm.train('1')` function.

Print the training and testing accuracy using `lm.predict` and `lm.compute_accuracy` given.

```

In [8]: lm=LogisticRegression()
lm.load_data('./data/logistic-regression-train.csv', './data/logistic-regression-test.csv')
training_accuracy= 0
testing_accuracy= 0
#####
# STRART YOUR CODE HERE #
#####
lm.normalize()
beta = lm.train('1')
train_predict = lm.predict(lm.train_x.values, beta)
test_predict = lm.predict(lm.test_x.values, beta)
training_accuracy = lm.compute_accuracy(train_predict, lm.train_y.values)
testing_accuracy = lm.compute_accuracy(test_predict, lm.test_y.values)
#####
# END YOUR CODE HERE #
#####
print('Training accuracy is: ', training_accuracy)
print('Testing accuracy is: ', testing_accuracy)

```

```

average logL for iteration 0: -0.7139731229909018
average logL for iteration 500: -0.5558529778896625
average logL for iteration 1000: -0.49990618171705625
average logL for iteration 1500: -0.4766183024021117
average logL for iteration 2000: -0.46678224234711174
average logL for iteration 2500: -0.4627302012028891
average logL for iteration 3000: -0.4611124973709076
average logL for iteration 3500: -0.46048365917256584
average logL for iteration 4000: -0.46024395049229505
average logL for iteration 4500: -0.4601537801568773
average logL for iteration 5000: -0.4601201520274992
average logL for iteration 5500: -0.46010767881536613
average logL for iteration 6000: -0.4601030679263344
average logL for iteration 6500: -0.46010136699582815
average logL for iteration 7000: -0.46010074033149045
average logL for iteration 7500: -0.46010050963239185
average logL for iteration 8000: -0.4601004247433812
average logL for iteration 8500: -0.4601003935162356
average logL for iteration 9000: -0.460100382031068
average logL for iteration 9500: -0.46010037780733604
Training avgLogL: -0.4601003762559442
Training accuracy is: 0.797
Testing accuracy is: 0.7534791252485089

```


Questions:

1. Compare the accuracy on the testing dataset for each version. Are they the same? Why or why not?
2. Regularization. Similar to linear regression, an regularization term could be added to logistic regression. The objective function becomes following:

$$J(\beta) = -\frac{1}{n} \sum_i (y_i x_i^T \beta - \log(1 + \exp\{x_i^T \beta\})) + \lambda \sum_j \beta_j^2,$$

where $\lambda \geq 0$, which is a hyper parameter that controls the trade off. Take the derivative $\frac{\partial J(\beta)}{\partial \beta_j}$ of this provided objective function and provide the batch gradient descent update.

Your answer here:

1. They are the same with a very small difference in their precisions. And it makes sense as they both have the same optimal point; both algorithms are guaranteed to converge.
- 2.

$$\beta_{new} = \beta_{old} + \frac{\partial J(\beta)}{\partial \beta_j}$$

$$\frac{\partial J(\beta)}{\partial \beta_j} = (XY - X\sigma(\beta) + 2\lambda\beta)$$

2.3 Visualize the decision boundary on a toy dataset

In this subsection, you will use the same implementation for another small dataset with each datapoint x with only two features (x_1, x_2) to visualize the decision boundary of logistic regression model.

```
In [9]: from hw1code.logistic_regression import LogisticRegression

lm=LogisticRegression(verbose = False)
lm.load_data('./data/logistic-regression-toy.csv', './data/logistic-regre
ssion-toy.csv')
# As a sanity check, we print out the size of the training data (99,2) a
nd training labels (99,)
print('Training data shape: ', lm.train_x.shape)
print('Training labels shape:', lm.train_y.shape)

Training data shape: (99, 2)
Training labels shape: (99,)
```

In the following block, you can apply the same implementation of logistic regression model (either in 2.1 or 2.2) to the toy dataset. Print out the $\hat{\beta}$ after training and accuracy on the train set.

```
In [10]: training_accuracy= 0
#=====#
#  STRART YOUR CODE HERE  #
#=====#
lm.normalize()
beta = lm.train('1')
train_predict = lm.predict(lm.train_x.values, beta)
training_accuracy = lm.compute_accuracy(train_predict,lm.train_y.values
)
#=====#
#    END YOUR CODE HERE    #
#=====#
print('Training accuracy is: ', training_accuracy)
print('Beta: ', beta)
```

```
Training avgLogL:  -0.3291474312957121
Training accuracy is:  0.8888888888888888
Beta:  [-0.04717577  1.46005896  2.06586134]
```

Next, we try to plot the decision boundary of your learned logistic regression classifier. Generally, a decision boundary is the region of a space in which the output label of a classifier is ambiguous. That is, in the given toy data, given a datapoint $x = (x_1, x_2)$ on the decision boundary, the logistic regression classifier cannot decide whether $y = 0$ or $y = 1$.

Question

Is the decision boundary for logistic regression linear? Why or why not?

Your answer here:

Yes, because the probability can be written as a linear function $p = \text{beta.dot}(x)$ which is linear as we are classifying the two features and the line separates them


Draw the decision boundary in the following cell. Note that the code to plot the raw data points are given. You may need `plt.plot` function (see [here](https://matplotlib.org/tutorials/introductory/pyplot.html) (<https://matplotlib.org/tutorials/introductory/pyplot.html>)).

```

In [14]: # scatter plot the raw data
df = pd.concat([lm.train_x, lm.train_y], axis=1)
groups = df.groupby("y")
for name, group in groups:
    plt.plot(group["x1"], group["x2"], marker="o", linestyle="", label=name)

# plot the decision boundary on top of the scattered points
#=====#
# STRART YOUR CODE HERE #
#=====#
db = -(beta[0] + beta[1] * df[['x1', 'x2']]) / beta[2]
plt.plot(X, db, 'r^')
#=====#
# END YOUR CODE HERE #
#=====#
plt.show()

```



End of Homework 1 :)

After you've finished the homework, please print out the entire `ipynb` notebook and two `py` files into one PDF file. Make sure you include the output of code cells and answers for questions. Prepare submit it to GradeScope.

```
1 import pandas as pd
2 import numpy as np
3 import sys
4 import random as rd
5
6 #insert an all-one column as the first column
7 def addAllOneColumn(matrix):
8     n = matrix.shape[0] #total of data points
9     p = matrix.shape[1] #total number of attributes
10
11     newMatrix = np.zeros((n,p+1))
12     newMatrix[:,1:] = matrix
13     newMatrix[:,0] = np.ones(n)
14
15     return newMatrix
16
17 # Reads the data from CSV files, converts it into Dataframe and returns x and
18 # y dataframes
19 def getDataframe(filePath):
20     dataframe = pd.read_csv(filePath)
21     y = dataframe['y']
22     x = dataframe.drop('y', axis=1)
23     return x, y
24
25 # train_x and train_y are numpy arrays
26 # function returns value of beta calculated using (0) the formula  $\beta = (X^T X)^{-1} (X^T Y)$ 
27 def getBeta(train_x, train_y):
28     n = train_x.shape[0] #total of data points
29     p = train_x.shape[1] #total number of attributes
30
31     beta = np.zeros(p)
32     #=====#
33     # STRART YOUR CODE HERE #
34     #=====#
35     # 1. Calculate the transpose
36     # 2. first_term =  $(X^T X)^{-1}$ 
37     # 3. second_term =  $(X^T Y)$ 
38     # 4. calculate beta = first_term * second_term
39     train_x_transpose = np.transpose(train_x)
40     first_term = np.linalg.inv(np.matmul(train_x_transpose, train_x))
41     second_term = np.matmul(train_x_transpose, train_y)
42     beta = np.matmul(first_term, second_term)
43     #=====#
44     # END YOUR CODE HERE #
45     #=====#
46     return beta
47
48 # train_x and train_y are numpy arrays
49 # lr (learning rate) is a scalar
50 # function returns value of beta calculated using (1) batch gradient descent
51 def getBetaBatchGradient(train_x, train_y, lr, num_iter):
52     beta = np.random.rand(train_x.shape[1])
53
54     n = train_x.shape[0] #total of data points
55     p = train_x.shape[1] #total number of attributes
56
57     beta = np.random.rand(p)
```

```

59     for iter in range(0, num_iter):
60         deriv = np.zeros(p)
61         for i in range(n):
62             #=====#
63             # STRART YOUR CODE HERE #
64             #=====#
65             x_i = train_x[i]
66             y_i = train_y[i]
67             x_i_transpose = np.transpose(x_i)
68             A = np.matmul(x_i_transpose, beta)
69             deriv += x_i * (A - y_i)
70             #=====#
71             #   END YOUR CODE HERE   #
72             #=====#
73         deriv = deriv / n
74         beta = beta - deriv.dot(lr)
75     return beta
76
77 # train_x and train_y are numpy arrays
78 # lr (learning rate) is a scalar
79 # function returns value of beta calculated using (2) stochastic gradient
80 # descent
81 def getBetaStochasticGradient(train_x, train_y, lr):
82     n = train_x.shape[0] #total of data points
83     p = train_x.shape[1] #total number of attributes
84
85     beta = np.random.rand(p)
86
87     epoch = 100;
88     for iter in range(epoch):
89         indices = list(range(n))
90         rd.shuffle(indices)
91         for i in range(n):
92             idx = indices[i]
93             #=====#
94             # STRART YOUR CODE HERE #
95             #=====#
96             y_i = train_y[idx]
97             x_i = train_x[idx]
98             x_i_transpose = np.transpose(x_i)
99             A = np.matmul(x_i_transpose, beta)
100            beta = beta + (lr * (y_i - A) * x_i)
101            #=====#
102            #   END YOUR CODE HERE   #
103            #=====#
104        return beta
105
106 # Linear Regression implementation
107 class LinearRegression(object):
108     # Initializes by reading data, setting hyper-parameters, and forming
109     # linear model
110     # Forms a linear model (learns the parameter) according to type of beta
111     # (0 - closed form, 1 - batch gradient, 2 - stochastic gradient)
112     # Performs z-score normalization if z_score is 1
113     def __init__(self, lr=0.005, num_iter=1000):
114         self.lr = lr
115         self.num_iter = num_iter
116         self.train_x = pd.DataFrame()

```

```

116     self.test_x = pd.DataFrame()
117     self.test_y = pd.DataFrame()
118     self.algType = 0
119     self.isNormalized = 0
120
121     def load_data(self, train_file, test_file):
122         self.train_x, self.train_y = getDataframe(train_file)
123         self.test_x, self.test_y = getDataframe(test_file)
124
125     def normalize(self):
126         # Applies z-score normalization to the dataframe and returns a
normalized dataframe
127         self.isNormalized = 1
128         means = self.train_x.mean(0)
129         std = self.train_x.std(0)
130         self.train_x = (self.train_x - means).div(std)
131         self.test_x = (self.test_x - means).div(std)
132
133         # Gets the beta according to input
134     def train(self, algType):
135         self.algType = algType
136         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one
column as the first column
137         print('Learning Algorithm Type: ', algType)
138         if(algType == '0'):
139             beta = getBeta(newTrain_x, self.train_y.values)
140             #print('Beta: ', beta)
141
142         elif(algType == '1'):
143             beta = getBetaBatchGradient(newTrain_x, self.train_y.values,
self.lr, self.num_iter)
144             #print('Beta: ', beta)
145         elif(algType == '2'):
146             self.lr = 0.0005 #lr to 0.0005 so that the beta does converge
147             beta = getBetaStochasticGradient(newTrain_x, self.train_y.values,
self.lr)
148             #print('Beta: ', beta)
149         else:
150             print('Incorrect beta_type! Usage: 0 - closed form solution, 1 -
batch gradient descent, 2 - stochastic gradient descent')
151
152         return beta
153
154     # Predicts the y values of all test points
155     # Outputs the predicted y values to the text file named "logistic-
regression-output_algType_isNormalized" inside "output" folder
156     # Computes MSE
157     def predict(self, x, beta):
158         newTest_x = addAllOneColumn(x)
159         self.predicted_y = newTest_x.dot(beta)
160         return self.predicted_y
161
162
163     # predicted_y and test_y are the predicted and actual y values
respectively as numpy arrays
164     # function prints the mean squared error value for the test dataset
165     def compute_mse(self, predicted_y, y):
166         mse = np.sum((predicted_y - y)**2)/predicted_y.shape[0]
167

```

169
170
171

```

1  #-*- coding: utf-8 -*-
2
3  import pandas as pd
4  import numpy as np
5  import sys
6  import random as rd
7
8  #insert an all-one column as the first column
9  def addAllOneColumn(matrix):
10     n = matrix.shape[0] #total of data points
11     p = matrix.shape[1] #total number of attributes
12
13     newMatrix = np.zeros((n,p+1))
14     newMatrix[:,0] = np.ones(n)
15     newMatrix[:,1:] = matrix
16
17
18     return newMatrix
19
20 # Reads the data from CSV files, converts it into Dataframe and returns x and
    y dataframes
21 def getDataframe(filePath):
22     dataframe = pd.read_csv(filePath)
23     y = dataframe['y']
24     x = dataframe.drop('y', axis=1)
25     return x, y
26
27 # sigmoid function
28 def sigmoid(z):
29     return 1 / (1 + np.exp(-z))
30
31 # compute average logL
32 def compute_avglogL(X,y,beta):
33     eps = 1e-50
34     n = y.shape[0]
35     avglogL = 0
36     #=====#
37     # STRART YOUR CODE HERE #
38     #=====#
39     for i in range(n):
40
41         x_transpose = np.transpose(X[i])
42         x_transpose_dot_beta = np.dot(x_transpose, beta)
43         first_term = y[i] * x_transpose_dot_beta
44         second_term = 1 + np.exp(x_transpose_dot_beta)
45         avglogL += first_term - np.log(second_term)
46
47     avglogL = avglogL/ n
48     #=====#
49     #   END YOUR CODE HERE   #
50     #=====#
51     return avglogL
52
53
54 # train_x and train_y are numpy arrays
55 # lr (learning rate) is a scalar
56 # function returns value of beta calculated using (0) batch gradient descent
57 def getBeta_BatchGradient(train_x, train_y, lr, num_iter, verbose):
58     beta = np.random.rand(train_x.shape[1])

```



```

60 n = train_x.shape[0] #total of data points
61 p = train_x.shape[1] #total number of attributes
62
63
64 beta = np.random.rand(p)
65 #update beta iteratively
66 for iter in range(0, num_iter):
67     #=====#
68     # STRART YOUR CODE HERE #
69     #=====#
70     for i in range(n):
71
72         beta_transpose_dot_x = np.dot(np.transpose(beta), train_x[i])
73         sigmoid_res = sigmoid(beta_transpose_dot_x)
74         diff = train_y[i] - sigmoid_res
75         gradient = np.dot(diff, train_x[i])
76         beta += gradient * lr
77     #=====#
78     #   END YOUR CODE HERE   #
79     #=====#
80     if(verbose == True and iter % 1000 == 0):
81         avgLogL = compute_avglogL(train_x, train_y, beta)
82         print(f'average logL for iteration {iter}: {avgLogL} \t')
83     return beta
84
85 # train_x and train_y are numpy arrays
86 # function returns value of beta calculated using (1) Newton-Raphson method
87 def getBeta_Newton(train_x, train_y, num_iter, verbose):
88     n = train_x.shape[0] #total of data points
89     p = train_x.shape[1] #total number of attributes
90
91     beta = np.random.rand(p)
92     ##### Please Fill Missing Lines Here #####
93     for iter in range(0, num_iter):
94         #=====#
95         # STRART YOUR CODE HERE #
96         #=====#
97         beta_XT = np.dot(beta, np.transpose(train_x))
98         sigmoid_res = sigmoid(beta_XT)
99         diff = train_y - sigmoid_res
100        # first deriv
101        first_deriv = np.dot(diff, train_x)
102        # second deriv
103        prob_mul = sigmoid_res * (1 - sigmoid_res)
104        x_mul = np.array([x*y for (x,y) in zip(train_x, prob_mul)])
105        second_deriv = -1 * np.dot(np.transpose(x_mul), train_x)
106        beta -= np.dot(np.linalg.inv(second_deriv), first_deriv)/n
107        #=====#
108        #   END YOUR CODE HERE   #
109        #=====#
110        if(verbose == True and iter % 500 == 0):
111            avgLogL = compute_avglogL(train_x, train_y, beta)
112            print(f'average logL for iteration {iter}: {avgLogL} \t')
113    return beta
114
115
116
117 # Linear Regression implementation
118 class LogisticRegression(object):

```

```

119     # Initializes by reading data, setting hyper-parameters, and forming
linear model
120     # Forms a linear model (learns the parameter) according to type of beta
(0 - batch gradient, 1 - Newton-Raphson)
121     # Performs z-score normalization if isNormalized is 1
122     # Print intermidate training loss if verbose = True
123     def __init__(self, lr=0.005, num_iter=10000, verbose = True):
124         self.lr = lr
125         self.num_iter = num_iter
126         self.verbose = verbose
127         self.train_x = pd.DataFrame()
128         self.train_y = pd.DataFrame()
129         self.test_x = pd.DataFrame()
130         self.test_y = pd.DataFrame()
131         self.algType = 0
132         self.isNormalized = 0
133
134
135     def load_data(self, train_file, test_file):
136         self.train_x, self.train_y = getDataframe(train_file)
137         self.test_x, self.test_y = getDataframe(test_file)
138
139     def normalize(self):
140         # Applies z-score normalization to the dataframe and returns a
normalized dataframe
141         self.isNormalized = 1
142         data = np.append(self.train_x, self.test_x, axis = 0)
143         means = data.mean(0)
144         std = data.std(0)
145         self.train_x = (self.train_x - means).div(std)
146         self.test_x = (self.test_x - means).div(std)
147
148     # Gets the beta according to input
149     def train(self, algType):
150         self.algType = algType
151         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one
column as the first column
152         if(algType == '0'):
153             beta = getBeta_BatchGradient(newTrain_x, self.train_y.values,
self.lr, self.num_iter, self.verbose)
154             #print('Beta: ', beta)
155
156             elif(algType == '1'):
157                 beta = getBeta_Newton(newTrain_x, self.train_y.values,
self.num_iter, self.verbose)
158                 #print('Beta: ', beta)
159             else:
160                 print('Incorrect beta_type! Usage: 0 - batch gradient descent, 1
- Newton-Raphson method')
161
162         train_avglogL = compute_avglogL(newTrain_x, self.train_y.values,
beta)
163         print('Training avgLogL: ', train_avglogL)
164
165         return beta
166
167     # Predicts the y values of all test points
168     # Outputs the predicted y values to the text file named "logistic-
regression-output_algType_isNormalized" inside "output" folder

```

```
170     def predict(self, x, beta):
171         newTest_x = addAllOneColumn(x)
172         self.predicted_y = (sigmoid(newTest_x.dot(beta))>=0.5)
173         return self.predicted_y
174
175     # predicted_y and y are the predicted and actual y values respectively as
numpy arrays
176     # function prints the accuracy
177     def compute_accuracy(self, predicted_y, y):
178         acc = np.sum(predicted_y == y)/predicted_y.shape[0]
179         return acc
180
```