

```

1  #-*- coding: utf-8 -*-
2
3  import pandas as pd
4  import numpy as np
5  import sys
6  import random as rd
7
8  #insert an all-one column as the first column
9  def addAllOneColumn(matrix):
10     n = matrix.shape[0] #total of data points
11     p = matrix.shape[1] #total number of attributes
12
13     newMatrix = np.zeros((n,p+1))
14     newMatrix[:,0] = np.ones(n)
15     newMatrix[:,1:] = matrix
16
17
18     return newMatrix
19
20 # Reads the data from CSV files, converts it into Dataframe and returns x and
    y dataframes
21 def getDataframe(filePath):
22     dataframe = pd.read_csv(filePath)
23     y = dataframe['y']
24     x = dataframe.drop('y', axis=1)
25     return x, y
26
27 # sigmoid function
28 def sigmoid(z):
29     return 1 / (1 + np.exp(-z))
30
31 # compute average logL
32 def compute_avglogL(X,y,beta):
33     eps = 1e-50
34     n = y.shape[0]
35     avglogL = 0
36     #=====#
37     # STRART YOUR CODE HERE #
38     #=====#
39     for i in range(n):
40
41         x_transpose = np.transpose(X[i])
42         x_transpose_dot_beta = np.dot(x_transpose, beta)
43         first_term = y[i] * x_transpose_dot_beta
44         second_term = 1 + np.exp(x_transpose_dot_beta)
45         avglogL += first_term - np.log(second_term)
46
47     avglogL = avglogL/ n
48     #=====#
49     #   END YOUR CODE HERE   #
50     #=====#
51     return avglogL
52
53
54 # train_x and train_y are numpy arrays
55 # lr (learning rate) is a scalar
56 # function returns value of beta calculated using (0) batch gradient descent
57 def getBeta_BatchGradient(train_x, train_y, lr, num_iter, verbose):
58     beta = np.random.rand(train_x.shape[1])

```

```

60 n = train_x.shape[0] #total of data points
61 p = train_x.shape[1] #total number of attributes
62
63
64 beta = np.random.rand(p)
65 #update beta iteratively
66 for iter in range(0, num_iter):
67     #=====#
68     # STRART YOUR CODE HERE #
69     #=====#
70     for i in range(n):
71
72         beta_transpose_dot_x = np.dot(np.transpose(beta), train_x[i])
73         sigmoid_res = sigmoid(beta_transpose_dot_x)
74         diff = train_y[i] - sigmoid_res
75         gradient = np.dot(diff, train_x[i])
76         beta += gradient * lr
77     #=====#
78     #   END YOUR CODE HERE   #
79     #=====#
80     if(verbose == True and iter % 1000 == 0):
81         avgLogL = compute_avglogL(train_x, train_y, beta)
82         print(f'average logL for iteration {iter}: {avgLogL} \t')
83     return beta
84
85 # train_x and train_y are numpy arrays
86 # function returns value of beta calculated using (1) Newton-Raphson method
87 def getBeta_Newton(train_x, train_y, num_iter, verbose):
88     n = train_x.shape[0] #total of data points
89     p = train_x.shape[1] #total number of attributes
90
91     beta = np.random.rand(p)
92     ##### Please Fill Missing Lines Here #####
93     for iter in range(0, num_iter):
94         #=====#
95         # STRART YOUR CODE HERE #
96         #=====#
97         beta_XT = np.dot(beta, np.transpose(train_x))
98         sigmoid_res = sigmoid(beta_XT)
99         diff = train_y - sigmoid_res
100        # first deriv
101        first_deriv = np.dot(diff, train_x)
102        # second deriv
103        prob_mul = sigmoid_res * (1 - sigmoid_res)
104        x_mul = np.array([x*y for (x,y) in zip(train_x, prob_mul)])
105        second_deriv = -1 * np.dot(np.transpose(x_mul), train_x)
106        beta -= np.dot(np.linalg.inv(second_deriv), first_deriv)/n
107        #=====#
108        #   END YOUR CODE HERE   #
109        #=====#
110        if(verbose == True and iter % 500 == 0):
111            avgLogL = compute_avglogL(train_x, train_y, beta)
112            print(f'average logL for iteration {iter}: {avgLogL} \t')
113    return beta
114
115
116
117 # Linear Regression implementation
118 class LogisticRegression(object):

```

```

119     # Initializes by reading data, setting hyper-parameters, and forming
linear model
120     # Forms a linear model (learns the parameter) according to type of beta
(0 - batch gradient, 1 - Newton-Raphson)
121     # Performs z-score normalization if isNormalized is 1
122     # Print intermidate training loss if verbose = True
123     def __init__(self, lr=0.005, num_iter=10000, verbose = True):
124         self.lr = lr
125         self.num_iter = num_iter
126         self.verbose = verbose
127         self.train_x = pd.DataFrame()
128         self.train_y = pd.DataFrame()
129         self.test_x = pd.DataFrame()
130         self.test_y = pd.DataFrame()
131         self.algType = 0
132         self.isNormalized = 0
133
134
135     def load_data(self, train_file, test_file):
136         self.train_x, self.train_y = getDataframe(train_file)
137         self.test_x, self.test_y = getDataframe(test_file)
138
139     def normalize(self):
140         # Applies z-score normalization to the dataframe and returns a
normalized dataframe
141         self.isNormalized = 1
142         data = np.append(self.train_x, self.test_x, axis = 0)
143         means = data.mean(0)
144         std = data.std(0)
145         self.train_x = (self.train_x - means).div(std)
146         self.test_x = (self.test_x - means).div(std)
147
148     # Gets the beta according to input
149     def train(self, algType):
150         self.algType = algType
151         newTrain_x = addAllOneColumn(self.train_x.values) #insert an all-one
column as the first column
152         if(algType == '0'):
153             beta = getBeta_BatchGradient(newTrain_x, self.train_y.values,
self.lr, self.num_iter, self.verbose)
154             #print('Beta: ', beta)
155
156             elif(algType == '1'):
157                 beta = getBeta_Newton(newTrain_x, self.train_y.values,
self.num_iter, self.verbose)
158                 #print('Beta: ', beta)
159             else:
160                 print('Incorrect beta_type! Usage: 0 - batch gradient descent, 1
- Newton-Raphson method')
161
162         train_avglogL = compute_avglogL(newTrain_x, self.train_y.values,
beta)
163         print('Training avgLogL: ', train_avglogL)
164
165         return beta
166
167     # Predicts the y values of all test points
168     # Outputs the predicted y values to the text file named "logistic-
regression-output_algType_isNormalized" inside "output" folder

```

```
170     def predict(self, x, beta):
171         newTest_x = addAllOneColumn(x)
172         self.predicted_y = (sigmoid(newTest_x.dot(beta))>=0.5)
173         return self.predicted_y
174
175     # predicted_y and y are the predicted and actual y values respectively as
numpy arrays
176     # function prints the accuracy
177     def compute_accuracy(self, predicted_y, y):
178         acc = np.sum(predicted_y == y)/predicted_y.shape[0]
179         return acc
180
```