

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 class TwoLayerNet(object):
5     """
6     A two-layer fully-connected neural network. The net has an input
7     dimension of
8     N, a hidden layer dimension of H, and performs classification over C
9     classes.
10    We train the network with a softmax loss function and L2 regularization
11    on the
12    weight matrices. The network uses a ReLU nonlinearity after the first
13    fully
14    connected layer.
15
16    In other words, the network has the following architecture:
17
18    input - fully connected layer - ReLU - fully connected layer - MSE Loss
19
20    ReLU function:
21    (i)  $x = x$  if  $x \geq 0$  (ii)  $x = 0$  if  $x < 0$ 
22
23    The outputs of the second fully-connected layer are the scores for each
24    class.
25    """
26
27    def __init__(self, input_size, hidden_size, output_size, std=1e-4):
28        """
29        Initialize the model. Weights are initialized to small random values
30        and
31        biases are initialized to zero. Weights and biases are stored in the
32        variable self.params, which is a dictionary with the following keys:
33
34        W1: First layer weights; has shape (H, D)
35        b1: First layer biases; has shape (H,)
36        W2: Second layer weights; has shape (C, H)
37        b2: Second layer biases; has shape (C,)
38
39        Inputs:
40        - input_size: The dimension D of the input data.
41        - hidden_size: The number of neurons H in the hidden layer.
42        - output_size: The number of classes C.
43        """
44        self.params = {}
45        self.params['W1'] = std * np.random.randn(hidden_size, input_size)
46        self.params['b1'] = np.zeros(hidden_size)
47        self.params['W2'] = std * np.random.randn(output_size, hidden_size)
48        self.params['b2'] = np.zeros(output_size)
49
50    def loss(self, X, y=None, reg=0.0):
51        """
52        Compute the loss and gradients for a two layer fully connected neural
53        network.
54
55        Inputs:
56        - X: Input data of shape (N, D). Each X[i] is a training sample.
57        - y: Vector of training labels. y[i] is the label for X[i], and each
58        y[i] is
59        an integer in the range  $0 \leq y[i] < C$ . This parameter is optional;
60        if it

```

```

53         is not passed then we only return scores, and if it is passed then
we
54         instead return the loss and gradients.
55         - reg: Regularization strength.
56
57     Returns:
58     If y is None, return a matrix scores of shape (N, C) where scores[i,
c] is
59     the score for class c on input X[i].
60
61     If y is not None, instead return a tuple of:
62     - loss: Loss (data loss and regularization loss) for this batch of
training
63     samples.
64     - grads: Dictionary mapping parameter names to gradients of those
parameters
65     with respect to the loss function; has the same keys as
self.params.
66     """
67     # Unpack variables from the params dictionary
68     W1, b1 = self.params['W1'], self.params['b1']
69     W2, b2 = self.params['W2'], self.params['b2']
70     N, D = X.shape
71
72     # Compute the forward pass
73     scores = None
74
75     # ===== #
76     # START YOUR CODE HERE
77     # ===== #
78     # Calculate the output scores of the neural network. The result
79     # should be (N, C). As stated in the description for this class,
80     # there should not be a ReLU layer after the second fully-connected
81     # layer.
82     # The code is partially given
83     # The output of the second fully connected layer is the output
scores.
84     # Do not use a for loop in your implementation.
85     # Please use 'h1' as input of hidden layers, and 'a2' as output of
86     # hidden layers after ReLU activation function.
87     # [Input X] --W1,b1--> [h1] -ReLU-> [a2] --W2,b2--> [scores]
88     # You may simply use np.maximum for implementing ReLU.
89     # Note that there is only one ReLU layer.
90     # Note that please do not change the variable names (h1, h2, a2)
91     # ===== #
92     h1 = np.dot(X, W1.T) + b1
93     a2 = np.zeros(h1.shape)
94     a2 = np.maximum(a2, h1)
95     h2 = np.dot(a2, W2.T) + b2
96     scores = h2
97     # ===== #
98     # END YOUR CODE HERE
99     # ===== #
100
101
102     # If the targets are not given then jump out, we're done
103     if y is None:
104         return scores
105
106     # Compute the loss

```

```

107     loss = None
108
109     # scores is num_examples by num_classes (N, C)
110     def softmax_loss(x, y):
111         loss, dx = 0,0
112         #
113         ===== #
114         # START YOUR CODE HERE (BONUS QUESTION)
115         #
116         ===== #
117         # Calculate the cross entropy loss after softmax output layer.
118         # The format are provided in the notebook.
119         # This function should return loss and dx, same as MSE loss
120         function.
121         #
122         ===== #
123         # END YOUR CODE HERE
124         #
125         ===== #
126         return loss, dx
127
128     def MSE_loss(x, y):
129         loss, dx = 0,0
130         #
131         ===== #
132         # START YOUR CODE HERE
133         #
134         ===== #
135         # This function should return loss and dx (gradients ready for
136         # back prop).
137         # The loss is MSE loss between network ouput and one hot vector
138         # of class
139         # labels is required for backpropogation.
140         #
141         ===== #
142         # Hint: Check the type and shape of x and y.
143         # e.g. print('DEBUG:x.shape, y.shape', x.shape, y.shape)
144         n = x.shape[0]
145         feature_size = x.shape[1]
146         target_matrix = np.zeros((n, feature_size))
147         for i in range(n):
148             j = y[i]
149             target_matrix[i][j] = 1
150             diff = x - target_matrix
151             loss = 0.5 * np.sum(np.square(diff)) / n
152             dx = diff / n
153             #
154             ===== #
155             # END YOUR CODE HERE
156             #
157             ===== #
158             return loss, dx
159
160     # data_loss, dscore = softmax_loss(scores, y)

```

```

154         # The above line is for bonus question. If you have implemented
softmax_loss, de-comment this line instead of MSE error.
155
156         data_loss, dscore = MSE_loss(scores, y) # "comment" this line if you
use softmax_loss
157         # ===== #
158         # START YOUR CODE HERE
159         # ===== #
160         # Calculate the regularization loss. Multiply the regularization
161         # loss by 0.5 (in addition to the factor reg).
162         # ===== #
163         reg_loss = 0.5 * reg * (np.sum(np.square(self.params['W1'])) +
np.sum(np.square(self.params['W2'])))
164         # ===== #
165         # END YOUR CODE HERE
166         # ===== #
167         loss = data_loss + reg_loss
168
169         grads = {}
170
171         # ===== #
172         # START YOUR CODE HERE
173         # ===== #
174         # Backpropagation: (You do not need to change this!)
175         # Backward pass is implemented. From the dscore error, we calculate
176         # the gradient and store as grads['W1'], etc.
177         # ===== #
178         grads['W2'] = a2.T.dot(dscore).T + reg * W2
179         grads['b2'] = np.ones(N).dot(dscore)
180
181         da_h = np.zeros(h1.shape)
182         da_h[h1>0] = 1
183         dh = (dscore.dot(W2) * da_h)
184
185         grads['W1'] = np.dot(dh.T,X) + reg * W1
186         grads['b1'] = np.ones(N).dot(dh)
187         # ===== #
188         # END YOUR CODE HERE
189         # ===== #
190
191         return loss, grads
192
193     def train(self, X, y, X_val, y_val,
194               learning_rate=1e-3, learning_rate_decay=0.95,
195               reg=1e-5, num_iters=100,
196               batch_size=200, verbose=False):
197         """
198         Train this neural network using stochastic gradient descent.
199
200         Inputs:
201         - X: A numpy array of shape (N, D) giving training data.
202         - y: A numpy array of shape (N,) giving training labels; y[i] = c
means that
203           X[i] has label c, where 0 <= c < C.
204         - X_val: A numpy array of shape (N_val, D) giving validation data.
205         - y_val: A numpy array of shape (N_val,) giving validation labels.
206         - learning_rate: Scalar giving learning rate for optimization.
207         - learning_rate_decay: Scalar giving factor used to decay the
learning rate
208           after each epoch.

```

```

209 - reg: Scalar giving regularization strength.
210 - num_iters: Number of steps to take when optimizing.
211 - batch_size: Number of training examples to use per step.
212 - verbose: boolean; if true print progress during optimization.
213 """
214 num_train = X.shape[0]
215 iterations_per_epoch = max(num_train / batch_size, 1)
216
217 # Use SGD to optimize the parameters in self.model
218 loss_history = []
219 train_acc_history = []
220 val_acc_history = []
221
222 for it in np.arange(num_iters):
223     X_batch = None
224     y_batch = None
225
226     # Create a minibatch (X_batch, y_batch) by sampling batch_size
227     # samples randomly.
228
229     b_index = np.random.choice(num_train, batch_size)
230     X_batch = X[b_index]
231     y_batch = y[b_index]
232
233     # Compute loss and gradients using the current minibatch
234     loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
235     loss_history.append(loss)
236
237     #
===== #
238     # START YOUR CODE HERE
239     #
===== #
240     # Perform a gradient descent step using the minibatch to update
241     # all parameters (i.e., W1, W2, b1, and b2).
242     # The gradient has been calculated as grads['W1'], grads['W2'],
243     # grads['b1'], grads['b2']
244     # For example,
245     # W1(new) = W1(old) - learning_rate * grads['W1']
246     # (this is not the exact code you use!)
247     #
===== #
248     self.params['b1'] = self.params['b1'] - learning_rate *
grads['b1']
249     self.params['b2'] = self.params['b2'] - learning_rate *
grads['b2']
250     self.params['W1'] = self.params['W1'] - learning_rate *
grads['W1']
251     self.params['W2'] = self.params['W2'] - learning_rate *
grads['W2']
252     #
===== #
253     # END YOUR CODE HERE
254     #
===== #
255
256     if verbose and it % 100 == 0:
257         print('iteration {} / {}: loss {}'.format(it, num_iters,
loss))
258

```

```

259         # Every epoch, check train and val accuracy and decay learning
rate.
260         if it % iterations_per_epoch == 0:
261             # Check accuracy
262             train_acc = (self.predict(X_batch) == y_batch).mean()
263             val_acc = (self.predict(X_val) == y_val).mean()
264             train_acc_history.append(train_acc)
265             val_acc_history.append(val_acc)
266
267             # Decay learning rate
268             learning_rate *= learning_rate_decay
269
270         return {
271             'loss_history': loss_history,
272             'train_acc_history': train_acc_history,
273             'val_acc_history': val_acc_history,
274         }
275
276     def predict(self, X):
277         """
278         Use the trained weights of this two-layer network to predict labels
for
279         data points. For each data point we predict scores for each of the C
280         classes, and assign each data point to the class with the highest
score.
281
282         Inputs:
283         - X: A numpy array of shape (N, D) giving N D-dimensional data points
to
284             classify.
285
286         Returns:
287         - y_pred: A numpy array of shape (N,) giving predicted labels for
each of
288             the elements of X. For all i, y_pred[i] = c means that X[i] is
predicted
289             to have class c, where 0 <= c < C.
290         """
291         y_pred = None
292
293         # ===== #
294         # START YOUR CODE HERE
295         # ===== #
296         # Predict the class given the input data.
297         # ===== #
298         scores = self.loss(X)
299         y_pred = np.argmax(scores, axis=1)
300         # ===== #
301         # END YOUR CODE HERE
302         # ===== #
303
304         return y_pred
305
306
307

```