

# CS145 Homework 6, Naive Bayes and Topic Modeling

**Due date:** HW6 is due on **11:59 PM PT, Dec. 14 (Monday, Final Week)**. Please submit through GradeScope.

---

## Print Out Your Name and UID

**Name:** Ali Mirabzadeh, **UID:** 305179067

---

## Important Notes about HW6

- HW6, as the last homework, is optional if you choose to use the first 5 homework assignments for homework grading. We will select your highest 5 homework grades to calculate your final homework grade.
  - Since HW6 is optional, for the implementation of Naive Bayes and pLSA, you can choose to implement the provided `.py` and `.py` file by filling in the blocks. **Alternatively, you are given the option to implement completely from scratch based on your understanding. Note that some packages with ready-to-use implementation of Naive Bayes and pLSA are not allowed.**
- 

## Before You Start

You need to first create HW6 conda environment by the given `cs145hw6.yml` file, which provides the name and necessary packages for this tasks. If you have `conda` properly installed, you may create, activate or deactivate by the following commands:

```
conda env create -f cs145hw6.yml
conda activate hw6
conda deactivate
```

OR

```
conda env create --name NAMEOFOURCHOICE -f cs145hw6.yml
conda activate NAMEOFOURCHOICE
conda deactivate
```

To view the list of your environments, use the following command:

```
conda env list
```

More useful information about managing environments can be found [here](https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html) (<https://docs.conda.io/projects/conda/en/latest/user-guide/tasks/manage-environments.html>).

You may also quickly review the usage of basic Python and Numpy package, if needed in coding for matrix operations.

In this notebook, you must not delete any code cells in this notebook. If you change any code outside the blocks (such as hyperparameters) that you are allowed to edit (between `STRART/END YOUR CODE HERE`), you need to highlight these changes. You may add some additional cells to help explain your results and observations.

```
In [1]: import numpy as np
from numpy import zeros, int8, log
from pylab import random
import pandas as pd
import matplotlib.pyplot as plt
from pylab import rcParams
rcParams['figure.figsize'] = 8,8
import seaborn as sns; sns.set()
import re
import time
import nltk
nltk.download('punkt')
from nltk.tokenize import word_tokenize
from sklearn.metrics import confusion_matrix

[nltk_data] Downloading package punkt to
[nltk_data]      /Users/alimirabzadeh/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
```

Note that `seaborn` in HW6 is only used for plotting classification confusion matrix (in a "heatmap" style). If you encounter installation problem and cannot solve it, you may use alternative plot methods to show your results.

## Section 1: Naive Bayes for Text (50 points)

Naive Bayes is one generative model for text classification. In the problem, you are given a document in `dataset` folder. The original data comes from ["20 newsgroups"](http://qwone.com/~jason/20Newsgroups/) (<http://qwone.com/~jason/20Newsgroups/>). You can use the provided data files to save efforts on preprocessing.

Note: The code and dataset are under the subfolder named `nb`.

```

In [2]: ### Data processing and preparation
# read train/test labels from files
train_label = pd.read_csv('./nb/dataset/train.label', names=['t'])
train_label = train_label['t'].tolist()
test_label = pd.read_csv('./nb/dataset/test.label', names=['t'])
test_label = test_label['t'].tolist()

# read train/test documents from files
train_data = open('./nb/dataset/train.data')
df_train = pd.read_csv(train_data, delimiter=' ', names=['docIdx', 'wordIdx'])
test_data = open('./nb/dataset/test.data')
df_test = pd.read_csv(test_data, delimiter=' ', names=['docIdx', 'wordIdx'])

# read vocab
vocab = open('./nb/dataset/vocabulary.txt')
vocab_df = pd.read_csv(vocab, names=['word'])
vocab_df = vocab_df.reset_index()
vocab_df['index'] = vocab_df['index'].apply(lambda x: x+1)

# add label column to original df_train
docIdx = df_train['docIdx'].values
i = 0
new_label = []
for index in range(len(docIdx)-1):
    new_label.append(train_label[i])
    if docIdx[index] != docIdx[index+1]:
        i += 1
new_label.append(train_label[i])
df_train['classIdx'] = new_label

```

If you have the data prepared properly, the following line of code would return the head of the `df_train` dataframe, which is,

	docIdx	wordIdx	count	classIdx
0	1	1	4	1
1	1	2	2	1
2	1	3	10	1
3	1	4	4	1
4	1	5	2	1

```

In [3]: # check the head of 'df_train'
print(df_train.head())

```

	docIdx	wordIdx	count	classIdx
0	1	1	4	1
1	1	2	2	1
2	1	3	10	1
3	1	4	4	1
4	1	5	2	1

Complete the implementation of Naive Bayes model for text classification `nbm.py`. After that, run `nbm_sklearn.py`, which uses `sklearn` to implement naive bayes model for text classification.

(Note that the dataset is slightly different loaded in `nbm_sklearn.py` and also you don't need to change anything in `nbm_sklearn.py` and directly run it.)

If the implementation is correct, you can expect the results are generally close on both train set accuracy and test set accuracy.

```
In [4]: from nb.nbm import NB_model

# model training
nbm = NB_model()
nbm.fit(df_train, train_label, vocab_df)
```

Prior Probability of each class:

```
1: 0.04259472890229834
2: 0.05155736977549028
3: 0.05075871860857219
4: 0.05208980388676901
5: 0.051024935664211554
6: 0.052533498979501284
7: 0.051646108794036735
8: 0.052533498979501284
9: 0.052888455053687104
10: 0.0527109770165942
11: 0.05306593309078002
12: 0.0527109770165942
13: 0.05244475996095483
14: 0.0527109770165942
15: 0.052622237998047744
16: 0.05315467210932647
17: 0.04836276510781791
18: 0.05004880646020055
19: 0.04117490460555506
20: 0.033365870973467035
```

Training completed!

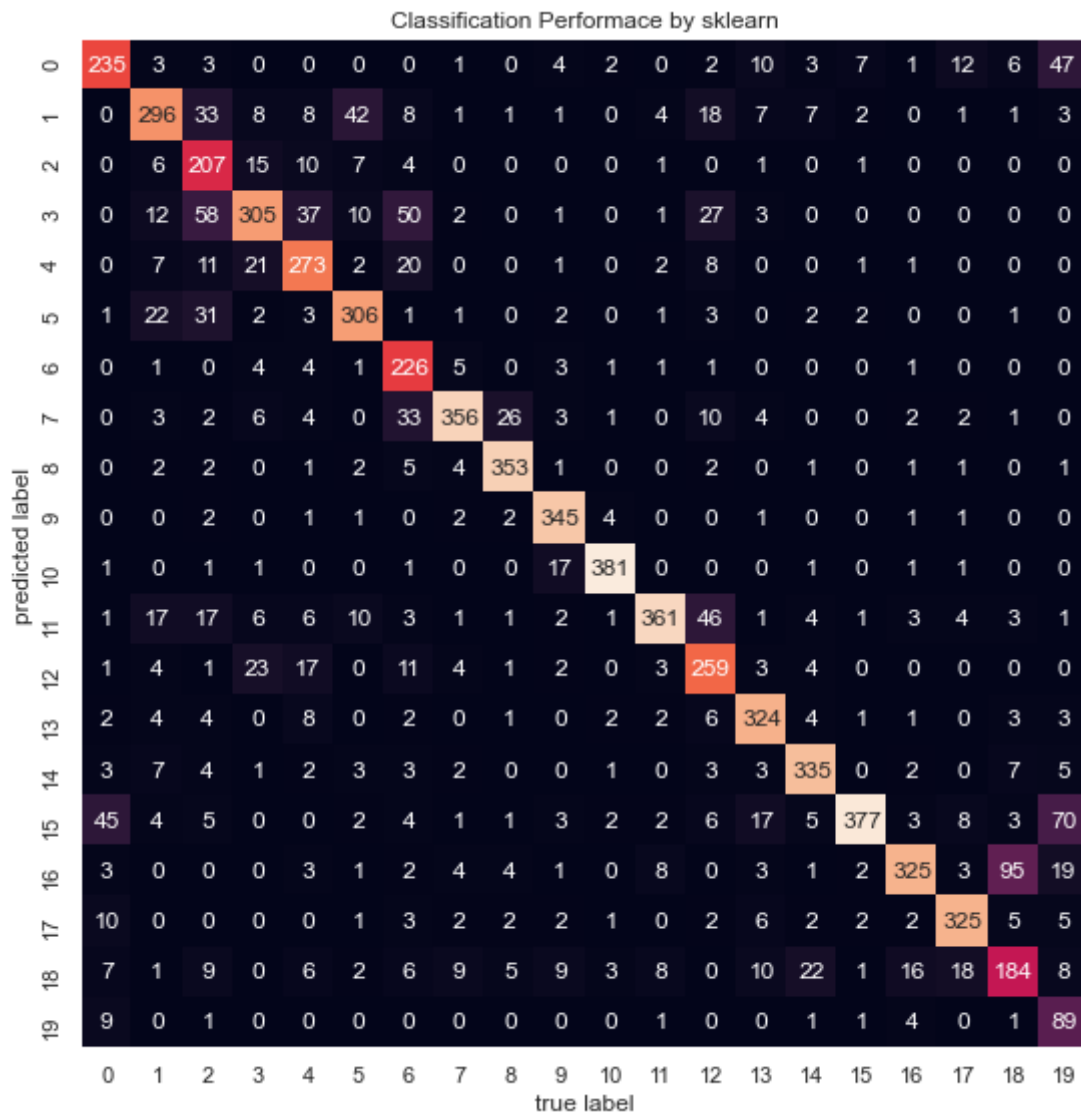
```
In [5]: # make predictions on train set to validate the model
predict_train_labels = nbm.predict(df_train)
train_acc = (np.array(train_label) == np.array(predict_train_labels)).mean()
print("Accuracy on training data by my implementation: {}".format(train_acc))

# make predictions on test data
predict_test_labels = nbm.predict(df_test)
test_acc = (np.array(test_label) == np.array(predict_test_labels)).mean()
print("Accuracy on training data by my implementation: {}".format(test_acc))
```

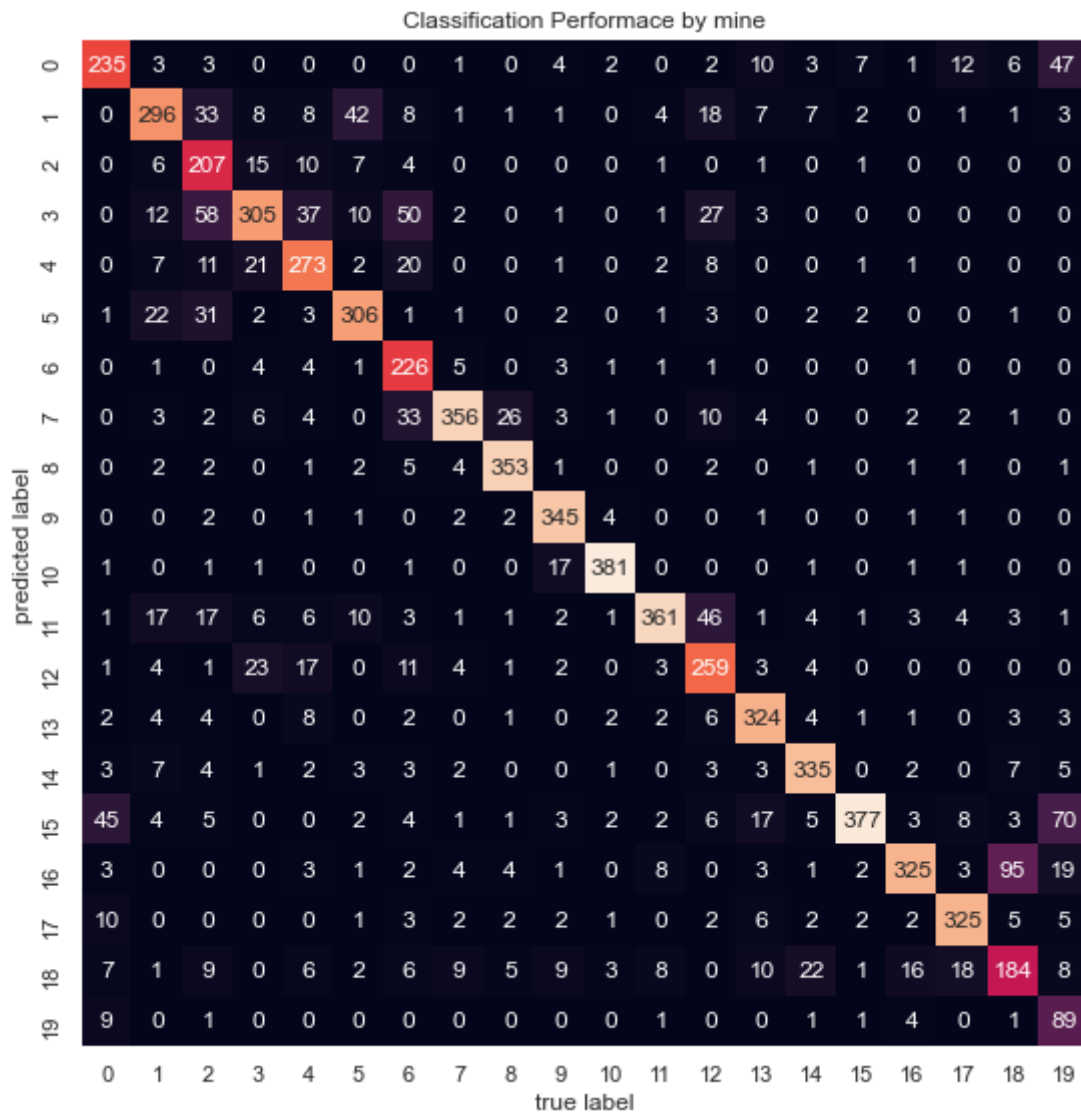
Accuracy on training data by my implementation: 0.941077291685154

Accuracy on training data by my implementation: 0.7810792804796802

```
In [6]: # plot classification matrix
mat = confusion_matrix(test_label, predict_test_labels)
sns.heatmap(mat.T, square=True, annot=True, fmt='d', cbar=False)
plt.title('Classification Performance by sklearn')
plt.xlabel('true label')
plt.ylabel('predicted label')
plt.tight_layout()
plt.savefig('./nb/output/nbm_sklearn.png')
plt.show()
```



```
In [7]: # plot classification matrix
mat = confusion_matrix(test_label, predict_test_labels)
sns.heatmap(mat.T, square=True, annot=True, fmt='d', cbar=False)
plt.title('Classification Performance by mine')
plt.xlabel('true label')
plt.ylabel('predicted label')
plt.tight_layout()
plt.savefig('./nb/output/nbm_mine.png')
plt.show()
##They seem to be identical!
```



**Reminder:** Do not forget to run nbm\_sklearn.py to compare the results to get the accuracy and confusion matrix by sklearn implementation. You can run `python nbm_sklearn.py` under the folder path of `./hw6/nb/`.

```

ions in the public API at pandas.testing instead.
import pandas.util.testing as tm
Accuracy on training data by sklearn: 0.9326498143892522
Accuracy on test data by sklearn: 0.7738980350504514
(here) -> sklearn.metrics.accuracy_score(y_test, y_pred)

```

## Question & Analysis

0. Please indicate whether you implemented based the given code or from scratch.
1. Report your classification accuracy on train and test documents. Also report your classification confusion matrix. Show one example document that Naive Bayes classifies incorrectly (i.e. fill in the following result table). Attach the output figure `./output/nbm_mine.png` in the jupyter book and briefly explain your observation on the accuracy and confusion matrix.

	Train set accuracy	Test set accuracy
sklearn implementaion	0.9326	0.7738
your implementaion	0.9411	0.7811

2. Show one example document that Naive Bayes classifies incorrectly by filling the following table. Provide your thought on the reason why this document is misclassified. (Note that the topic mapping is available at `train.map` same as `test.map`)

Words (count) in the example document	Predicted label	Truth label
sit (2), couple (1), ...	Class 20	Class 3

3. Is Naive Bayes a generative model or discriminative model and why? What is the difference between Naive Bayes classifier and Logistic Regression? What are the pros and cons of Naive Bayes for text classification task?
4. Can you apply Naive Bayes model to identify spam emails from normal ones? Briefly explain your method (you don't need to implementation for this question).

## Your Answers

### 0. Given Code

1. My implementation's results are pretty close to the one from Sklearn as we can in the above example. I think the reason for that is "Note that the dataset is slightly different loaded in `nbm_sklearn.py`". Also, the confusion matrixes seem identical
2. I basically used `train.data` and picked `docID = 11269` and picked a few word such as `sit` and `couple` and looked for their corresponding `docID` in `test.data`. Then I mapped the `docID` to the corresponding label using `train.label` and `test.label` and noticed for those word they were predicted class 20 even though they are class 3
3. It's a generative model because it learns from joint probability distribution. NB assumes that each feature is conditionally independent where as logistic regression doesn't make the same assumption and in fact it uses conditional probabilities.  
PROS: It's relatively simple to implement and is useful for applications like identifying spam/non-spam emails. CONS: However, since as mentioned, the model makes assumption that the features are independent so could misclassify as well that's why it's called naive.

4. Yes, in fact it's really great for indentifying spam from normal emails/message. Like we can train a model that has both spam and non-spam emails, so the model can learn what words appear in spams find the probabilities of each word

## Section 2: Topic Modeling: Probabilistic Latent Semantic Analysis (50 points)

In this section, you will implement Probabilistic Latent Semantic Analysis (pLSA) by EM algorithm. Note: The code and dataset are under the subfolder named `plsa`. You can find two dataset files named `dataset1.txt` and `dataset2.txt` together with a [stopword](https://en.wikipedia.org/wiki/Stop_word) ([https://en.wikipedia.org/wiki/Stop\\_word](https://en.wikipedia.org/wiki/Stop_word)) list as `stopwords.dic`.

First complete the implementation of pLSA in `plsa.py`. You need to finish the E step, M step and likelihood function. Note that the optimizing process on dataset 2 might take a while.

```
In [20]: # input file, output files and parameters
datasetFilePath = './plsa/dataset/dataset2.txt' # or set as './plsa/dataset
stopwordsFilePath = './plsa/dataset/stopwords.dic'
docTopicDist = './plsa/output/docTopicDistribution.txt'
topicWordDist = './plsa/output/topicWordDistribution.txt'
dictionary = './plsa/output/dictionary.dic'
topicWords = './plsa/output/topics.txt'

K = 10 # number of topic
maxIteration = 20 # maxIteration and threshold control the train process
threshold = 3
topicWordsNum = 10 # parameter for output
```

```
In [21]: from plsa.plsa import PLSA
from plsa.utils import preprocessing

N, M, word2id, id2word, X = preprocessing(datasetFilePath, stopwordsFilePat
```



```

In [22]: plsa_model = PLSA()
plsa_model.initialize(N, K, M, word2id, id2word, X)

oldLoglikelihood = 1
newLoglikelihood = 1
print ("K: ", K)
for i in range(0, maxIteration):
    plsa_model.EStep() #implement E step
    plsa_model.MStep() #implement M step
    newLoglikelihood = plsa_model.LogLikelihood()
    print("[",time.strftime('%Y-%m-%d %H:%M:%S',time.localtime(time.time()))
          "iteration", str(newLoglikelihood))
    # you should see increasing loglikelihood
    #if(newLoglikelihood - oldLoglikelihood < threshold):
        #break
    oldLoglikelihood = newLoglikelihood

plsa_model.output(docTopicDist, topicWordDist, dictionary, topicWords, topicWords)

```

```

K: 10
[ 2020-12-12 20:45:48 ] 1 iteration -152813.14222663164
[ 2020-12-12 20:46:38 ] 2 iteration -150702.2084671577
[ 2020-12-12 20:47:27 ] 3 iteration -147769.08482924715
[ 2020-12-12 20:48:15 ] 4 iteration -144286.93004346848
[ 2020-12-12 20:49:03 ] 5 iteration -140920.01085965018
[ 2020-12-12 20:49:51 ] 6 iteration -138065.39935262286
[ 2020-12-12 20:50:39 ] 7 iteration -135761.93322620235
[ 2020-12-12 20:51:28 ] 8 iteration -133966.65665138207
[ 2020-12-12 20:52:17 ] 9 iteration -132604.4753048485
[ 2020-12-12 20:53:04 ] 10 iteration -131577.26133772268
[ 2020-12-12 20:53:52 ] 11 iteration -130798.00244629064
[ 2020-12-12 20:54:41 ] 12 iteration -130206.31420406947
[ 2020-12-12 20:55:30 ] 13 iteration -129751.51912562801
[ 2020-12-12 20:56:18 ] 14 iteration -129396.17958538682
[ 2020-12-12 20:57:07 ] 15 iteration -129115.51666607056
[ 2020-12-12 20:57:55 ] 16 iteration -128892.35183639737
[ 2020-12-12 20:58:44 ] 17 iteration -128713.02264192027
[ 2020-12-12 20:59:33 ] 18 iteration -128565.99063177603
[ 2020-12-12 21:00:21 ] 19 iteration -128461.21019689328
[ 2020-12-12 21:01:12 ] 20 iteration -128391.02602565885

```

```

In [8]: plsa_model.output(docTopicDist, topicWordDist, dictionary, topicWords, topicWords)

```

K trials for dataset1:

**K: 2**

```
[ 2020-12-12 19:52:11 ] 1 iteration -7919.6262395904005
[ 2020-12-12 19:52:12 ] 2 iteration -7849.457857168658
[ 2020-12-12 19:52:12 ] 3 iteration -7748.688566278395
[ 2020-12-12 19:52:12 ] 4 iteration -7649.049330932226
[ 2020-12-12 19:52:12 ] 5 iteration -7570.229444533049
[ 2020-12-12 19:52:13 ] 6 iteration -7515.863292371991
[ 2020-12-12 19:52:13 ] 7 iteration -7486.754431830214
[ 2020-12-12 19:52:13 ] 8 iteration -7468.153382467457
[ 2020-12-12 19:52:13 ] 9 iteration -7453.440627442476
[ 2020-12-12 19:52:13 ] 10 iteration -7439.133104330273
[ 2020-12-12 19:52:14 ] 11 iteration -7423.579265547843
[ 2020-12-12 19:52:14 ] 12 iteration -7410.424259313294
[ 2020-12-12 19:52:14 ] 13 iteration -7403.122683031999
[ 2020-12-12 19:52:14 ] 14 iteration -7399.443413341776
[ 2020-12-12 19:52:14 ] 15 iteration -7397.904025162465
[ 2020-12-12 19:52:15 ] 16 iteration -7396.827965428653
[ 2020-12-12 19:52:15 ] 17 iteration -7395.167380502538
[ 2020-12-12 19:52:15 ] 18 iteration -7394.080746167426
[ 2020-12-12 19:52:15 ] 19 iteration -7393.566915161791
[ 2020-12-12 19:52:16 ] 20 iteration -7393.219834120998
```

**K: 3**

```
[ 2020-12-12 19:53:46 ] 1 iteration -7901.390411485891
[ 2020-12-12 19:53:46 ] 2 iteration -7800.116753824728
[ 2020-12-12 19:53:46 ] 3 iteration -7656.417216100979
[ 2020-12-12 19:53:46 ] 4 iteration -7482.062455212207
[ 2020-12-12 19:53:47 ] 5 iteration -7310.070139050328
[ 2020-12-12 19:53:47 ] 6 iteration -7180.512572669283
[ 2020-12-12 19:53:47 ] 7 iteration -7095.806817014209
[ 2020-12-12 19:53:48 ] 8 iteration -7046.064177397884
[ 2020-12-12 19:53:48 ] 9 iteration -7019.3430791795945
[ 2020-12-12 19:53:48 ] 10 iteration -6996.0751705925295
[ 2020-12-12 19:53:49 ] 11 iteration -6970.001732099258
[ 2020-12-12 19:53:49 ] 12 iteration -6944.556252137331
[ 2020-12-12 19:53:49 ] 13 iteration -6929.911450845423
[ 2020-12-12 19:53:49 ] 14 iteration -6924.797974358775
[ 2020-12-12 19:53:50 ] 15 iteration -6922.550293576286
[ 2020-12-12 19:53:50 ] 16 iteration -6921.153524156658
[ 2020-12-12 19:53:50 ] 17 iteration -6920.682085754938
[ 2020-12-12 19:53:51 ] 18 iteration -6920.508744484476
[ 2020-12-12 19:53:51 ] 19 iteration -6920.419385760949
[ 2020-12-12 19:53:51 ] 20 iteration -6920.367391100283
```

```

K: 4
[ 2020-12-12 19:28:13 ] 1 iteration -7741.817334532188
[ 2020-12-12 19:28:13 ] 2 iteration -7567.8917657387165
[ 2020-12-12 19:28:14 ] 3 iteration -7399.0793349558535
[ 2020-12-12 19:28:14 ] 4 iteration -7249.489181636675
[ 2020-12-12 19:28:14 ] 5 iteration -7119.101651314833
[ 2020-12-12 19:28:15 ] 6 iteration -7014.786765565352
[ 2020-12-12 19:28:15 ] 7 iteration -6925.490809474547
[ 2020-12-12 19:28:15 ] 8 iteration -6842.5491207282475
[ 2020-12-12 19:28:16 ] 9 iteration -6773.174608455707
[ 2020-12-12 19:28:16 ] 10 iteration -6728.555918522748
[ 2020-12-12 19:28:16 ] 11 iteration -6704.522138456633
[ 2020-12-12 19:28:17 ] 12 iteration -6691.289441504023
[ 2020-12-12 19:28:17 ] 13 iteration -6684.534950399586
[ 2020-12-12 19:28:18 ] 14 iteration -6680.0189138696
[ 2020-12-12 19:28:18 ] 15 iteration -6676.102057629343
[ 2020-12-12 19:28:18 ] 16 iteration -6671.675016285936
[ 2020-12-12 19:28:19 ] 17 iteration -6666.545915827314
[ 2020-12-12 19:28:19 ] 18 iteration -6662.238776656173
[ 2020-12-12 19:28:19 ] 19 iteration -6660.300068320686
[ 2020-12-12 19:28:20 ] 20 iteration -6659.607678141125

```

K trials for datasets2:

```

K: 2
[ 2020-12-12 19:55:54 ] 1 iteration -153741.30982181325
[ 2020-12-12 19:56:05 ] 2 iteration -152830.88672493157
[ 2020-12-12 19:56:16 ] 3 iteration -151977.3835246241
[ 2020-12-12 19:56:27 ] 4 iteration -151229.49484554326
[ 2020-12-12 19:56:38 ] 5 iteration -150643.34143026551
[ 2020-12-12 19:56:50 ] 6 iteration -150211.7626220788
[ 2020-12-12 19:57:01 ] 7 iteration -149901.37494689875
[ 2020-12-12 19:57:12 ] 8 iteration -149679.8382707446
[ 2020-12-12 19:57:23 ] 9 iteration -149518.03058236607
[ 2020-12-12 19:57:34 ] 10 iteration -149396.424813981
[ 2020-12-12 19:57:46 ] 11 iteration -149300.89956035835
[ 2020-12-12 19:57:57 ] 12 iteration -149220.96021128865
[ 2020-12-12 19:58:08 ] 13 iteration -149151.43672967708
[ 2020-12-12 19:58:19 ] 14 iteration -149088.02518471918
[ 2020-12-12 19:58:30 ] 15 iteration -149031.8590827375
[ 2020-12-12 19:58:42 ] 16 iteration -148985.96587131688
[ 2020-12-12 19:58:53 ] 17 iteration -148949.2021389379
[ 2020-12-12 19:59:04 ] 18 iteration -148920.43603557497
[ 2020-12-12 19:59:15 ] 19 iteration -148896.71806280757
[ 2020-12-12 19:59:26 ] 20 iteration -148874.20020718026

```

```

K: 3
[ 2020-12-12 20:10:31 ] 1 iteration -153507.93675527006
[ 2020-12-12 20:10:46 ] 2 iteration -152283.65440721923
[ 2020-12-12 20:11:01 ] 3 iteration -150904.24638959861
[ 2020-12-12 20:11:17 ] 4 iteration -149541.38862566577
[ 2020-12-12 20:11:33 ] 5 iteration -148365.17278460204
[ 2020-12-12 20:11:49 ] 6 iteration -147414.2755639266
[ 2020-12-12 20:12:04 ] 7 iteration -146666.59358327917
[ 2020-12-12 20:12:20 ] 8 iteration -146086.9223074089
[ 2020-12-12 20:12:36 ] 9 iteration -145646.452292398
[ 2020-12-12 20:12:51 ] 10 iteration -145323.4856398951
[ 2020-12-12 20:13:06 ] 11 iteration -145095.21431161382
[ 2020-12-12 20:13:22 ] 12 iteration -144933.12441075433
[ 2020-12-12 20:13:38 ] 13 iteration -144811.18897993598
[ 2020-12-12 20:13:54 ] 14 iteration -144712.65284691955
[ 2020-12-12 20:14:09 ] 15 iteration -144636.4382910918
[ 2020-12-12 20:14:25 ] 16 iteration -144580.80182003425
[ 2020-12-12 20:14:40 ] 17 iteration -144540.0237754329
[ 2020-12-12 20:14:56 ] 18 iteration -144509.36156574343
[ 2020-12-12 20:15:11 ] 19 iteration -144481.25919719363
[ 2020-12-12 20:15:26 ] 20 iteration -144454.04203414303

```

---

```

K: 4
[ 2020-12-12 20:23:51 ] 1 iteration -153150.16363130286
[ 2020-12-12 20:24:10 ] 2 iteration -151574.13582451956
[ 2020-12-12 20:24:30 ] 3 iteration -149662.70289579206
[ 2020-12-12 20:24:50 ] 4 iteration -147632.2332598618
[ 2020-12-12 20:25:10 ] 5 iteration -145848.33577260995
[ 2020-12-12 20:25:30 ] 6 iteration -144492.17210519378
[ 2020-12-12 20:25:50 ] 7 iteration -143533.16818947604
[ 2020-12-12 20:26:09 ] 8 iteration -142855.4735421343
[ 2020-12-12 20:26:29 ] 9 iteration -142355.54309563778
[ 2020-12-12 20:26:50 ] 10 iteration -141979.8650948237
[ 2020-12-12 20:27:11 ] 11 iteration -141699.89628784102
[ 2020-12-12 20:27:30 ] 12 iteration -141487.4629628496
[ 2020-12-12 20:27:51 ] 13 iteration -141322.9215173641
[ 2020-12-12 20:28:10 ] 14 iteration -141196.0950263997
[ 2020-12-12 20:28:31 ] 15 iteration -141094.0518121649
[ 2020-12-12 20:28:51 ] 16 iteration -141007.1258854785
[ 2020-12-12 20:29:11 ] 17 iteration -140934.09464703494
[ 2020-12-12 20:29:31 ] 18 iteration -140874.1458193671
[ 2020-12-12 20:29:52 ] 19 iteration -140827.51683027574
[ 2020-12-12 20:30:13 ] 20 iteration -140791.882043118

```

### Question & Analysis

0. Please indicate whether you implemented based the given code or from scratch.
1. Choose different  $K$  (number of topics) in `p1sa.py`. What is your option for a reasonable  $K$  in `dataset1.txt` and `dataset2.txt`? Give your results of 10 words under each topic by filling in the following table (suppose you set  $K = 4$ ).

For dataset 1:

Topic 1	Topic 2	Topic 3	Topic 4
luffy devil pirates fruit piece "" manga user fruits	luffy crew alabasta baroque navy ace d. pirates pirate roger	grand sea haki called island burū mountain pose blue red	luffy pirates crew island straw dressrosa franky alliance zou hats

For dataset 2:

Topic 1	Topic 2	Topic 3	Topic 4
"" percent soviet u.s. officials oil rate monday prices	"" bank percent soviet people gorbachev government billion economy	"" u.s. official government people president noriega roberts united	"" bush dukakis people campaign percent president company california

2. Are there any similarities between pLSA and GMM model? Briefly explain your thoughts.
3. What are the disadvantages of pLSA? Consider its generalizing ability to new unseen document and its parameter complexity, etc.

## Your Answers

### 0. Given Code

1. Based on the Piazza post I tried different K's from {2,3,4} and checking the topics results, k=4 seem to be reasonable
2. Yes, they both use EM algorithm and probability distribution. PLSA is applicable for text data.
3. I think tuning the hyperparameters is a disadvantage as for large datasets, like dataset2, here, it could take hours to find the best parameters. pLSA is not good at generalizing new data so this could results in not performing well on new unseen data

## Bonus Questions (10 points): LDA

We've learned document and topic modeling techniques. As mentioned in the lecture, most frequently used topic models are pLSA and LDA. [Latent Dirichlet allocation \(LDA\)](https://ai.stanford.edu/~ang/papers/nips01-lda) (<https://ai.stanford.edu/~ang/papers/nips01-lda>) proposed by David M. Blei, Andrew Y. Ng, and Michael I. Jordan, posits that each document is generated as a mixture of topics where the continuous-valued mixture proportions are distributed as a latent Dirichlet random variable.

In this question, please read the paper and/or tutorials of LDA and finish the following questions and tasks:

- (1) What are the differences between pLSA and LDA? List at least one advantage of LDA over pLSA?
- (2) Show a demo of LDA with brief result analysis on any corpus and discuss what real-world applications can be supported by LDA. Note: You do not need to implement LDA algorithms from scratch. You may use multiple packages such as `nltk`, `gensim`, `pyLDAvis` (added on the `cs145hw6.yml`) to help show the demo within couple of lines of code. If you'd like to use other packages, feel free to install them.

Your Answers Used: <https://medium.com/nanonets/topic-modeling-with-lsa-psla-lda-and->

lda2vec-

[555ff65b0b05#:~:text=LDA%20typically%20works%20better%20than,fixed%20point%20in%20the](#)

<https://medium.com/nanonets/topic-modeling-with-lsa-psla-lda-and-lda2vec->

[555ff65b0b05#:~:text=LDA%20typically%20works%20better%20than,fixed%20point%20in%20the](#)

1. "LDA typically works better than pLSA because it can generalize to new documents easily." As mentioned pLSA is not good at generalizing new unseen data so that's an advantage of LDA over pLSA

```
In [ ]: import nltk
import gensim
```

**End of Homework 6 :)**

Please printout the Jupyter notebook and relevant code files that you work on and submit only 1 PDF file on GradeScope with page assigned.

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import collections
5
6 class NB_model():
7     def __init__(self):
8         self.pi = {} # to store prior probability of each class
9         self.Pr_dict = None
10        self.num_vocab = None
11        self.num_classes = None
12
13    def fit(self, train_data, train_label, vocab, if_use_smooth=True):
14        # get prior probabilities
15        self.num_vocab = len(vocab['index'].tolist())
16        self.get_prior_prob(train_label)
17        # ===== YOUR CODE HERE =====
18        # Calculate probability of each word based on class
19        # Hint: Store each probability value in matrix or dict:
20        self.Pr_dict[classID][wordID] or Pr_dict[wordID][classID])
21        # Remember that there are possible NaN or 0 in Pr_dict matrix/dict.
22        Use smooth method
23        self.classes = collections.defaultdict(int)
24
25        word_count_per_class = collections.defaultdict(lambda:
26collections.defaultdict(int))
27        self.Pr_dict = collections.defaultdict(lambda:
28collections.defaultdict(float))
29
30        train_dict = train_data.to_dict()
31        for i in range(len(train_dict['classIdx'])):
32            self.classes[train_dict['classIdx'][i]] += train_dict['count'][i]
33            word_count_per_class[train_dict['classIdx'][i]]
34            [train_dict['wordIdx'][i]] += train_dict['count'][i]
35
36        for classID in word_count_per_class:
37            for wordID in word_count_per_class[classID]:
38                self.Pr_dict[classID][wordID] = (word_count_per_class[classID]
39[wordID] + 1) /
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999

```

```

53     score_dict = {}
54     max_score = 0
55     #Creating a probability row for each class
56     for classIdx in range(1,self.num_classes+1):
57         score_dict[classIdx] = 0
58         # ===== YOUR CODE HERE =====
59         ### Implement the score_dict for all classes for each document
60         ### Remember to use log addition rather than probability
multiplication
61         ### Remember to add prior probability, i.e. self.pi
62         score_dict[classIdx] += np.log(self.pi[classIdx])
63         for wordId in new_dict[docIdx]:
64             if self.Pr_dict[classIdx][wordIdx] == 0:
65                 score_dict[classIdx] += new_dict[docIdx][wordId] *
np.log(1/(self.classes[classIdx] + self.num_vocab))
66             else:
67                 score_dict[classIdx] += new_dict[docIdx][wordId] *
np.log(self.Pr_dict[classIdx][wordId])
68         # =====
69         max_score = max(score_dict, key=score_dict.get)
70         prediction.append(max_score)
71     return prediction
72
73
74     def get_prior_prob(self,train_label, verbose=True):
75         unique_class = list(set(train_label))
76         self.num_classes = len(unique_class)
77         total = len(train_label)
78         for c in unique_class:
79             # ===== YOUR CODE HERE =====
80             ### calculate prior probability of each class ###
81             ### Hint: store prior probability of each class in self.pi
82             counter = 0
83             for label in train_label:
84                 if c is label:
85                     counter += 1
86             self.pi[c] = counter / total
87             # =====
88         if verbose:
89             print("Prior Probability of each class:")
90             print("\n".join("{}: {}".format(k, v) for k, v in
self.pi.items()))
91

```



```

1 from numpy import zeros, int8, log
2 from pylab import random
3 import sys
4 #import jieba
5 import nltk
6 from nltk.tokenize import word_tokenize
7 import re
8 import time
9 import codecs
10 # N is # of document
11 # K is # of topic
12 # M is # of word
13 # beta is probability of word given a topic
14 # theta is probability of a topic given a document
15 # document- word matrix, N x M : word count in a document
16 class PLSA(object):
17     def initialize(self, N, K, M, word2id, id2word, X):
18         self.word2id, self.id2word, self.X = word2id, id2word, X
19         self.N, self.K, self.M = N, K, M
20         # theta[i, j] : p(zj|di): 2-D matrix
21         self.theta = random([N, K])
22         # beta[i, j] : p(wj|zi): 2-D matrix
23         self.beta = random([K, M])
24         # p[i, j, k] : p(zk|di,wj): 3-D tensor
25         self.p = zeros([N, M, K])
26         for i in range(0, N):
27             normalization = sum(self.theta[i, :])
28             for j in range(0, K):
29                 self.theta[i, j] /= normalization;
30
31         for i in range(0, K):
32             normalization = sum(self.beta[i, :])
33             for j in range(0, M):
34                 self.beta[i, j] /= normalization;
35
36
37     def EStep(self):
38         for i in range(0, self.N):
39             for j in range(0, self.M):
40                 ## ===== YOUR CODE HERE
41                 ### for each word in each document, calculate its
42                 ### conditional probability belonging to each topic (update
43                 p)
44                 denominator = 0
45                 for k in range(0, self.K):
46                     self.p[i, j, k] = self.theta[i, k] * self.beta[k, j]
47                     denominator += self.p[i, j, k]
48                 for k in range(0, self.K):
49                     self.p[i, j, k] /= denominator
50                 #
51                 =====
52
53     def MStep(self):
54         # update beta
55         for k in range(0, self.K):
56             # ===== YOUR CODE HERE =====
57             ### Implement M step 1: given the conditional distribution
58             ### find the parameters that can maximize the expected

```

```

57     denominator = 0
58     for m in range(0, self.M):
59         self.beta[k, m] = 0
60         for n in range(0, self.N):
61             self.beta[k, m] += self.X[n, m] * self.p[n, m, k]
62         denominator += self.beta[k, m]
63     for m in range(0, self.M):
64         self.beta[k, m] /= denominator
65     # =====
66
67     # update theta
68     for i in range(0, self.N):
69         # ===== YOUR CODE HERE =====
70         ### Implement M step 2: given the conditional distribution
71         ### find the parameters that can maximize the expected
likelihood (update theta)
72         for k in range(0, self.K):
73             self.theta[i, k] = 0
74             denominator = 0
75             for m in range(0, self.M):
76                 self.theta[i, k] += self.X[i, m] * self.p[i, m, k]
77             denominator += self.X[i, m]
78             self.theta[i, k] /= denominator
79         # =====
80
81
82     # calculate the log likelihood
83     def LogLikelihood(self):
84         loglikelihood = 0
85         for i in range(0, self.N):
86             for j in range(0, self.M):
87                 # ===== YOUR CODE HERE
=====
88                 ### Calculate likelihood function
89                 temp = 0
90                 for k in range(0, self.K):
91                     temp += self.theta[i, k] * self.beta[k, j]
92                 if temp > 0:
93                     loglikelihood += self.X[i, j] * log(second_term)
94                 #
=====
95         return loglikelihood
96
97     # output the params of model and top words of topics to files
98     def output(self, docTopicDist, topicWordDist, dictionary, topicWords,
topicWordsNum):
99         # document-topic distribution
100         file = codecs.open(docTopicDist, 'w', 'utf-8')
101         for i in range(0, self.N):
102             tmp = ''
103             for j in range(0, self.K):
104                 tmp += str(self.theta[i, j]) + ' '
105             file.write(tmp + '\n')
106         file.close()
107
108         # topic-word distribution
109         file = codecs.open(topicWordDist, 'w', 'utf-8')
110         for i in range(0, self.K):
111             tmp = ''

```

```
113         tmp += str(self.beta[i, j]) + ' '
114     file.write(tmp + '\n')
115 file.close()
116
117 # dictionary
118 file = codecs.open(dictionary, 'w', 'utf-8')
119 for i in range(0, self.M):
120     file.write(self.id2word[i] + '\n')
121 file.close()
122
123 # top words of each topic
124 file = codecs.open(topicWords, 'w', 'utf-8')
125 for i in range(0, self.K):
126     topicword = []
127     ids = self.beta[i, :].argsort()
128     for j in ids:
129         topicword.insert(0, self.id2word[j])
130     tmp = ''
131     for word in topicword[0:min(topicWordsNum, len(topicword))]:
132         tmp += word + ' '
133     file.write(tmp + '\n')
134 file.close()
```